




Photogrammetry & Robotics Lab

Machine Learning for Robotics and Computer Vision Tutorial

More on Regression

Jens Behley

Topics of this week lecture



- Linear Regression (ML, MAP, Bayesian)
- Closed-form solutions! (not always possible)

Recap: Linear Regression

- Under this assumptions, this leads to the following probabilistic formulation

$$P(y|\mathbf{x}, \theta) = \mathcal{N}(y|f(\mathbf{x}), \sigma^2)$$

- In linear regression, we assume that parameters θ appear **linearly** in our model

$$f(\mathbf{x}) = \mathbf{x}^T \theta + \theta_0$$

- θ_0 is called **intercept** (or **bias**) that enables us to have also functions that do not pass through the origin

Relation to Least Squares

- Maximum Likelihood:

$$\theta^* = \arg \min_{\theta} -\log \prod_{n=1}^N P(y_n | \mathbf{x}_n, \theta)$$

- Showed that results in NLL:

$$\mathcal{L}(\theta) := \frac{1}{2\sigma^2} \sum_{n=1}^N \underbrace{(y_n - \mathbf{x}_n^T \theta)^2}_{\text{L2 Loss}}$$

- NLL of Linear Regression is just Least Squares!
- (Losses are often denoted by $\ell(y_n, \hat{y}_n)$)

Recap: Non-linear Functions

- Linear regression is *linear in parameters*
- We can apply non-linear transformation:

$$f(\mathbf{x}) = \mathbf{x}^T \theta \longrightarrow f(\mathbf{x}) = \phi(\mathbf{x})^T \theta$$

- Let $\phi : \mathbb{R}^D \rightarrow \mathbb{R}^K$ and define $\Phi \in \mathbb{R}^{N \times K}$ as

$$\Phi = \begin{pmatrix} \phi(\mathbf{x}_1)^T \\ \vdots \\ \phi(\mathbf{x}_N)^T \end{pmatrix} \in \mathbb{R}^{N \times K}$$

- Everything else stays the same! Use normal equation with Φ instead of \mathbf{X}

$$\theta = (\Phi^T \Phi)^{-1} \Phi^T \mathbf{y}$$


Example: Polynomial transformation

- With polynomial transformation, we can fit polynomials of degree K

$$\phi_{\text{poly}}(\mathbf{x}_n) = \begin{pmatrix} 1 \\ x_1 \\ x_1^2 \\ \vdots \\ x_1^K \\ \vdots \end{pmatrix} \in \mathbb{R}^{DK+1}$$

- With K=1 it's “vanilla” linear regression

Recap: Example: Polynomial Fit



- With a polynomial of degree 3, we get a good fit. But can we do better with higher degrees?

Potential Problem

$$\phi_{\text{poly}}(\mathbf{x}_n) = \begin{pmatrix} 1 \\ x_1 \\ x_1^2 \\ \vdots \\ x_1^K \\ \vdots \end{pmatrix} \in \mathbb{R}^{DK+1}$$

- What happens when $\mathbf{x} = (10, 103, 1005)$ and we want $K = 10$?

Potential Problem

$$\phi_{\text{poly}}(\mathbf{x}_n) = \begin{pmatrix} 1 \\ x_1 \\ x_1^2 \\ \vdots \\ x_1^K \\ \vdots \end{pmatrix} \in \mathbb{R}^{DK+1}$$


- What happens when $\mathbf{x} = (10, 103, 1005)$ and we want $K = 10$?
 - $1005^{10} =$
10.511.401.320.40.790.642.597.666.015.625

**How can we avoid
numerical overflow?**


Solution: Feature Normalization

- Feature normalization ensures that polynomial transform get's not too large
 - **Important:** Same normalization to test samples
- Different options:
 - Divide by maximum value
 $x = (10, 103, 1005) / (40, 150, 2000)$
 $\rightarrow (0.25, 0.69, 0.50)$ $0.69^{10} = 0.0244$
 - Standardization (especially with image data)
 - Divide by mean and variance of training data
- If you provide pre-trained models: provide normalization constants


Recap: Example: Overfitting



degree 3



degree 5



degree 8

- Increasing the degree, we will get better training error
- But function will have implausible shape

Common Parameter Estimation

- Learning is finding parameters of

$$P(\theta|\mathbf{x}_{1:N}, y_{1:N}) \propto P(y_{1:N}|\mathbf{x}_{1:N}, \theta) \boxed{P(\theta)}$$

- Paradigms for parameter estimation:
 1. Point estimate with uniform prior
→ Maximum Likelihood Estimation
 2. Point estimate with given prior
→ Maximum A posteriori Estimation (MAP)
 3. Determine posterior over the parameters
→ Bayesian Estimation

NLL with Prior

- Assume Gaussian prior for parameters:

$$P(\theta) = \mathcal{N}(\theta|0, b^2 \mathbf{Id})$$

- NLL is then

$$\mathcal{L}_{\text{MAP}}(\theta) = -\log \prod_{n=1}^N P(y_n | \mathbf{x}_n, \theta) - \log P(\theta)$$

- Inserting Gaussians results in

Squared Length



$$\begin{aligned} \mathcal{L}_{MAP}(\theta) &= \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \phi(\mathbf{x}_n)^T \theta)^2 + \frac{1}{2b^2} \theta^T \theta + \text{const} \\ &= \frac{1}{2\sigma^2} (\mathbf{y} - \mathbf{\Phi}\theta)^T (\mathbf{y} - \mathbf{\Phi}\theta) + \frac{1}{2b^2} \theta^T \theta + \text{const} \end{aligned}$$

NLL with Prior

- Assume Gaussian prior for parameters:

$$P(\theta) = \mathcal{N}(\theta|0, b^2 \mathbf{Id})$$

- NLL is then

$$\mathcal{L}_{\text{MAP}}(\theta) = -\log \prod_{n=1}^N P(y_n | \mathbf{x}_n, \theta) - \log P(\theta)$$

- Inserting Gaussians results in

Squared Length




$$\mathcal{L}_{\text{MAP}}(\theta) = \frac{1}{2\sigma^2} \sum_{n=1}^N (y_n - \phi(\mathbf{x}_n)^T \theta)^2 + \frac{1}{2b^2} \theta^T \theta + \text{const}$$

$$= \frac{1}{2\sigma^2} (\mathbf{y} - \Phi\theta)^T (\mathbf{y} - \Phi\theta) + \boxed{\lambda} \theta^T \theta + \text{const}$$


Regularizer

Example: ML vs. MAP Estimate


— ML Estimate — MAP Estimate



degree 3




degree 5



degree 8

- Smoother functions even at higher degrees!

But there is a problem ...



$$\lambda = 0.1$$

- Let's say y is in $[1000, 1200]$

Any idea?

Bias term

$$= \frac{1}{2\sigma^2} (\mathbf{y} - \Phi\theta)^T (\mathbf{y} - \Phi\theta) + \boxed{\lambda} \theta^T \theta + \text{const}$$

- As also the bias term is regularized, but needs to be large...
- **Solution:** Exclude bias term in regularization (set $(\theta^T \theta)_0 = 0$)

See you next week!