

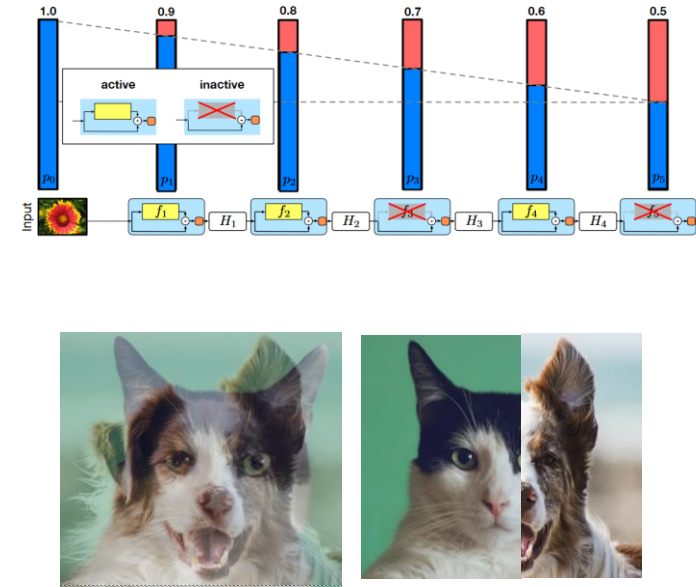
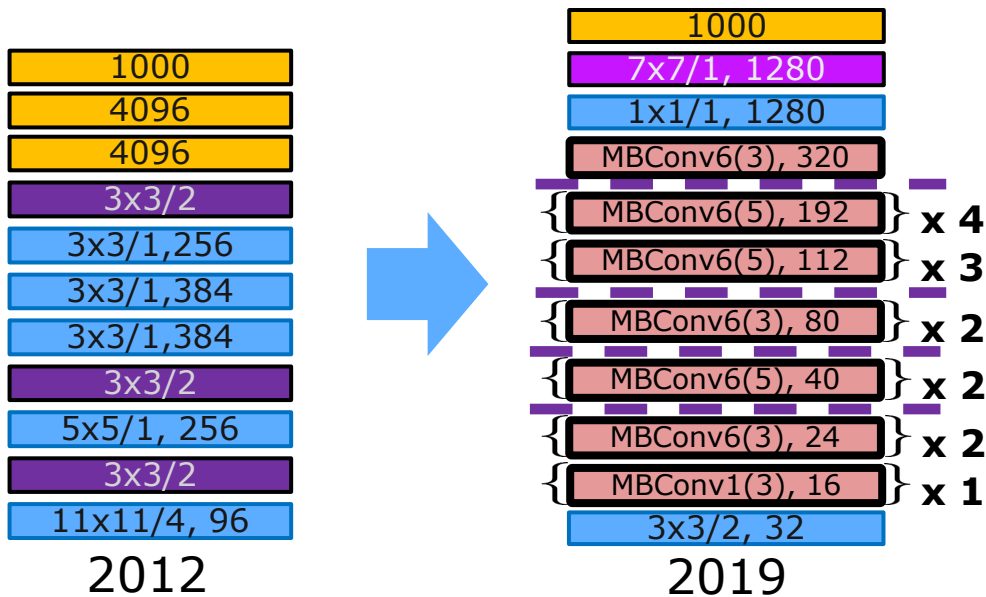
Photogrammetry & Robotics Lab

Machine Learning for Robotics and Computer Vision Tutorial

CNNs in Practice


Jens Behley

Last Lecture



- Popular and significant architectures and changes
 - global avg pooling → skip connection → efficiency
 - VGG → GoogleLeNet → ResNet → MobileNetV2 → EfficientNet
- Looked at common ways to close the gap between training and test performance (generalization gap)

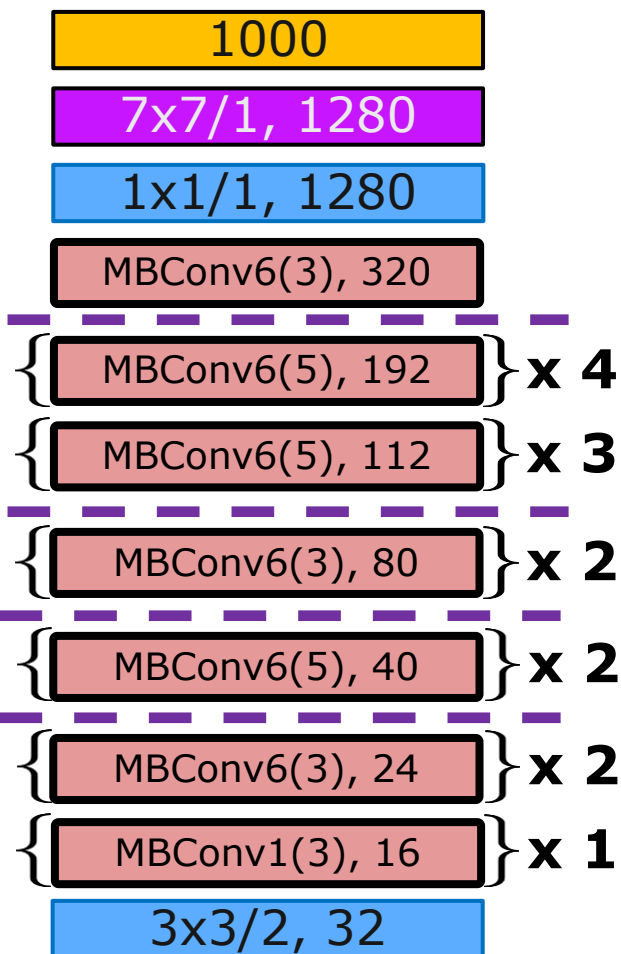
Recap: ResNet (2015)



- 18/34/50/101/152 learnable layers
- Downsampling (— —) via strided (s=2) convolution in first convolution of 2nd, 3rd, 4th stage
- ResNet-152 reaches 3.6% Top 5-error by ensemble of 6 models

 Convolution + ReLU
 Avg Pooling
 Max Pooling
 FC layer

Recap: EfficientNet (2019)



EfficientNet-B0

- NAS with objective of accuracy and computational efficiency (FLOPS)
- Structural very similar to MobileNetV2
 - Different number of channels
 - Convs with 3x3 and 5x5 kernels
- MBConv6(k) is inverted bottleneck with $t=6$ and $k \times k$ depth-wise separable convolutions

Model	Top-1	#Params
EfficientNet-B0	77.1	5.3M
ResNet-50	76.0	26 M
MobileNetV2	72.0	3.4M

Which architecture to choose?

- **ResNet-50** is a good baseline model that is often used and shows good performance
- Nowadays, **EfficientNets** are efficient alternative and current research focuses on more on training efficiency
- Vibrant research field → novel insights
- Improved training strategies shows promising improvement even for “old” ResNets

ResNet in PyTorch

```
import torchvision.models as models

# initialize resnet with FC layer
# outputting 10 scores instead of 1000
resnet50 = models.resnet50(num_classes=10)
```

- In torchvision are already many models implemented
- Rich set of options, see documentation.
- Can be used as building block (replace fc with identity module if fully-connected layer not needed)




Best Practices

- Some advice on how to approach a new problem (data, research problem, ...)
- If common vision task: get a quick baseline working
 - pre-trained networks, existing approaches, e.g., torchvision has off-the-shelf implementations for classification, detection, semantic/instance segmentation & tutorials
- Insights about the task; this is the thing that you want to “beat”

Be paranoid!

- Check everything, plot/visualize images
- Often subtle errors in parts that you think are easy, you did 1000 times
→ Check especially the easy parts
- Beauty, but also danger of Machine Learning: Even if the objective does not make sense or is erroneous, gradient descent will still optimize this.
- **Beware:** CNNs can still do a good job, even though there are errors in the pipeline or data

Example: Data augmentation in semantic segmentation



- Horizontal flip of image, must be also applied to pixel-wise labels
- A CNN can still produce reasonable results with only flipped images, but performance will drop!

Again: Getting the Data Right

- You will run many, many experiments, variations of your model, etc.
- Have a **validation set** ready that is large enough and can be used to assess your progress
- You want to know if you are making progress; if validation set changes all the time, you get “noisy” feedback and cannot compare to earlier steps

General workflow

- Common receipt
 1. First get **simple pipeline** working with small number of images
→ general problems or issues will not be magically solved by putting more data in!

General workflow

- Common receipt
 1. First get **simple pipeline** working with small number of images
→ general problems or issues will not be magically solved by putting more data in!
 2. **Monitor learning progress** via loss curves & metrics
→ Use tensorboard to visualize logs

General workflow

- Common receipt

1. First get **simple pipeline** working with small number of images
→ general problems or issues will not be magically solved by putting more data in!
2. **Monitor learning progress** via loss curves & metrics
→ Use tensorboard to visualize logs
3. **Improve model** by analyzing behavior on train/validation or dev set
→ adding additional things (maybe going **back to step 1**)

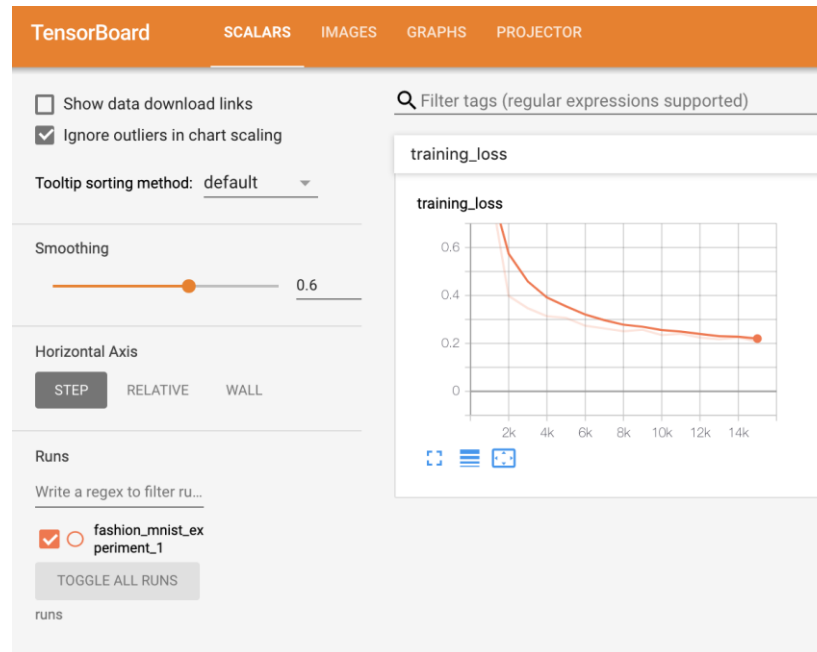
Step 1: Getting Started

- Always start with a simple pipeline and only some examples
- First experiments with only little data
- Does everything work as expected

Step 1: General Receipt


- Receipt:
 1. No data augmentation, regularization, dropout, etc. (added later)
 2. Single batch of few images (5-10 minibatches) → Network should be able to reach 100%
 3. Check if possible to overfit on small data
 4. Determine learning rate on all data on small number of iterations (LR: 0.1, 0.01, 0.001, ...)
 5. Determine better learning rate/weight decay (1-5 epochs)
 6. Train longer → Monitor progress

Step 2: Babysitting




- Use tensorboard to visualize loss (train, validation) & metrics of interest
- Visualizes logs that are produced while training
- See tutorial Pytorch's tutorial on tensorboard

Learning Curves



- Optimal curves look like decreasing fast, no larger “upticks”

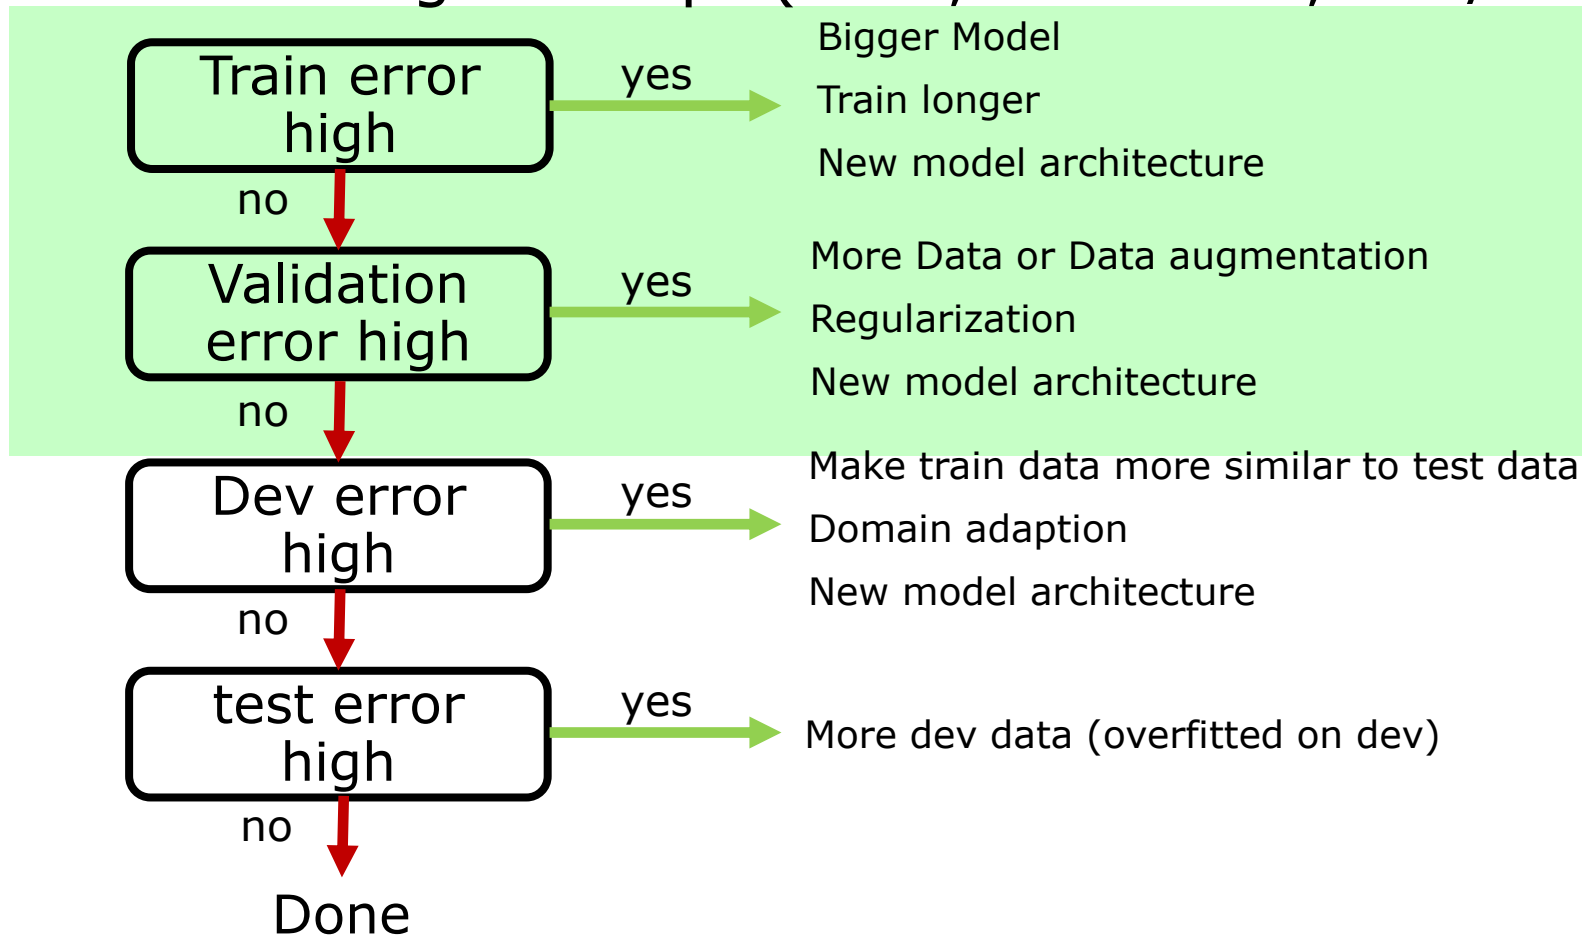
Learning Curves – LR too high



- Spikes are never a good sign, even if loss decreases afterwards
- Usually an indicator that LR is too high: decay LR before

Step 3: Improve Model

- Assess performance and improve model
- Andrew Ng's receipt (train, validation, dev/test set)



Step 3: Trying Things

- Make always one change, like adding data augmentation, adding layers, etc.
- Here, the fixed validation set will help to figure out if you are on the right track
- Get inspiration from similar work. Look at latest papers on ImageNet classification
- There is no single way or algorithm that you

Additional Reading

- Other resources on best practices & strategies:
 1. Justin Johnson's Notes in Lecture [PDF](#)
 2. Andre Karpathy's Notes [PDF](#)
 3. Andrew Ng. Machine Learning Yearning [PDF](#)

See you next week!